

A Preliminary Investigation into Storageless Blockchain

Yang Ji

April 15, 2019

1. Introduction
2. Analysis of Existing Works
3. Stateless Blockchain

Introduction

- In cryptocurrency, peer-to-peer payment transactions are asynchronously broadcasted and recorded in an ordered ledger.

Background

- In cryptocurrency, peer-to-peer payment transactions are asynchronously broadcasted and recorded in an ordered ledger.
- Consensus protocol requires nodes to validate new transactions.

- In cryptocurrency, peer-to-peer payment transactions are asynchronously broadcasted and recorded in an ordered ledger.
- Consensus protocol requires nodes to validate new transactions.
- e.g. To send 6 ETC from Alice to Bob requires that Alice has at least 6 ETC.

- In cryptocurrency, peer-to-peer payment transactions are asynchronously broadcasted and recorded in an ordered ledger.
- Consensus protocol requires nodes to validate new transactions.
- e.g. To send 6 ETC from Alice to Bob requires that Alice has at least 6 ETC.
- Simply querying the history from adjacent nodes is infeasible and insecure due to the large size of blocks.

Background

- In cryptocurrency, peer-to-peer payment transactions are asynchronously broadcasted and recorded in an ordered ledger.
- Consensus protocol requires nodes to validate new transactions.
- e.g. To send 6 ETC from Alice to Bob requires that Alice has at least 6 ETC.
- Simply querying the history from adjacent nodes is infeasible and insecure due to the large size of blocks.
- For that reason, most cryptocurrency nodes need to locally maintain the **validation state**, which means downloading all past transactions/accounts.

Background

- In Bitcoin, Zcash and Komodo, **validation state** is a set of immutable coins called **UTXO** (unspent transaction output)

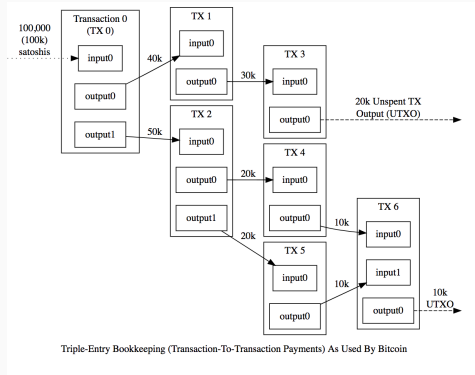


Fig. 1 : UTXO Model

Background

- In Ethereum, Nxt and Bitshares organize **validation state** as a set of mutable (and potential long-living) **accounts**.

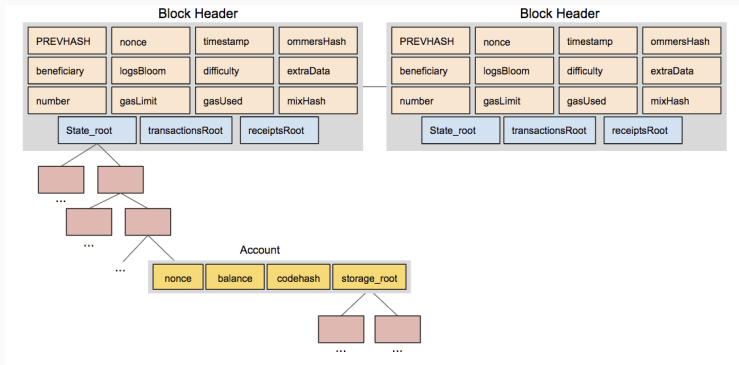


Fig. 2 : Account-based Model

- Other cryptocurrencies

System	Execution State	Proof Size	Bootstrap Security
Bitcoin [21]	UTXOs	Headers + TXs	Probabilistic (heaviest chain wins)
Ethereum [9]	All accounts	Headers + All accounts	Probabilistic (heaviest chain wins)
Permissioned	<u>Live accounts</u> Shards	Majority of trust set's signatures	Cryptographic if majority never compromised; none otherwise
OmniLedger [17] + Chainiac [22]	<u>UTXOs</u> Shards	$\frac{\text{Headers} + \text{Certificates}}{\text{Sparseness}} + \frac{\text{UTXOs}}{\text{Shards}}$	Cryptographic with static attacker; none with adaptive attacker
Algorand [14]	UTXOs	Headers + Certificates + TXs	Cryptographic
Vault	<u>Live accounts</u> Shards	$\frac{\text{Headers} + \text{Certificates}}{\text{Sparseness}} + \frac{\text{Live accounts}}{\text{Shards}}$	Cryptographic

Fig. 3 : Comparison of different cryptocurrency models

Motivation

- Maintaining ledger state is cumbersome from the perspectives of **storage** and **bootstrapping**
 - Large size (Bitcoin is around 150 GB and Ethereum has exceeded 400 GB)
 - Data storage size is linear with block number and could grow substantially in the coming years
 - Slow disk I/O operations (LevelDB or RocksDB)
 - DoS attack (adversarially-crafted transactions that need massive of disk accesses)
 - Increase the possibility of centralization in blockchains.
- Several people including Peter Todd and Mike Hearn also talked about **storageless clients** for Bitcoin in 2013.
- It thus appears that we need to compress the **ledger history** or allow miners to validate pending transactions without storing all past blocks.

Difficulty

- Currently the maximum size of the UTXO/Account set is unbounded as there is no consensus rule that limits growth.
- State set growth is driven by a number of factors, including the following fact.
 - For **Bitcoin**, there are massive of merge inputs, lost coins and dust outputs.
 - For **Ethereum**, smart contract inadvertently created many **zero-balance accounts** (account for around 38%).
- Based on these, a naïve way is to prune and clean 'useless' dust coins/accounts.
- However, there is little incentive to carry it out for many reasons:
 - Dust coins can't be economically spent and have other use cases.
 - We can't delete zero accounts because nodes need to keep track of the sequence number ("nonce") [Woo+].
- Redesigning index structures and storage mode has limited improvement.

Analysis of Existing Works

Possible Solutions

- Method 1: Storage Rent

Possible Solutions

- Method 1: **Storage Rent**
 - Miners could know and update ledger state by outsourcing local database to third party or sharing common files with other miners in IPFS system.

Possible Solutions

- Method 1: **Storage Rent**
 - Miners could know and update ledger state by outsourcing local database to third party or sharing common files with other miners in IPFS system.
 - **Drawbacks:** The operation cost is very high; less competitive; low I/O operations.

Possible Solutions

- Method 1: **Storage Rent**
 - Miners could know and update ledger state by outsourcing local database to third party or sharing common files with other miners in IPFS system.
 - **Drawbacks:** The operation cost is very high; less competitive; low I/O operations.
- Method 2: **Sharding**

Possible Solutions

- Method 1: **Storage Rent**
 - Miners could know and update ledger state by outsourcing local database to third party or sharing common files with other miners in IPFS system.
 - **Drawbacks:** The operation cost is very high; less competitive; low I/O operations.
- Method 2: **Sharding**
 - Partition the whole blockchain/storage into many shardings like Zilliqa [Tea+17] and Omnidex [KJG+18].

Possible Solutions

- Method 1: **Storage Rent**
 - Miners could know and update ledger state by outsourcing local database to third party or sharing common files with other miners in IPFS system.
 - **Drawbacks:** The operation cost is very high; less competitive; low I/O operations.
- Method 2: **Sharding**
 - Partition the whole blockchain/storage into many shardings like Zilliqa [Tea+17] and Omniledger [KJG+18].
 - **Drawbacks:** Security issues; Incentive mechanism;

Possible Solutions

- Method 1: **Storage Rent**
 - Miners could know and update ledger state by outsourcing local database to third party or sharing common files with other miners in IPFS system.
 - **Drawbacks:** The operation cost is very high; less competitive; low I/O operations.
- Method 2: **Sharding**
 - Partition the whole blockchain/storage into many shardings like Zilliqa [Tea+17] and Omniledger [KJG+18].
 - **Drawbacks:** Security issues; Incentive mechanism;
- Method 3: **Merging Transactions**

Possible Solutions

- Method 1: **Storage Rent**
 - Miners could know and update ledger state by outsourcing local database to third party or sharing common files with other miners in IPFS system.
 - **Drawbacks:** The operation cost is very high; less competitive; low I/O operations.
- Method 2: **Sharding**
 - Partition the whole blockchain/storage into many shardings like Zilliqa [Tea+17] and Omniledger [KJG+18].
 - **Drawbacks:** Security issues; Incentive mechanism;
- Method 3: **Merging Transactions**
 - Record fewer transactions using payment channels e.g. lightning network [PD16]

Possible Solutions

- Method 1: **Storage Rent**
 - Miners could know and update ledger state by outsourcing local database to third party or sharing common files with other miners in IPFS system.
 - **Drawbacks:** The operation cost is very high; less competitive; low I/O operations.
- Method 2: **Sharding**
 - Partition the whole blockchain/storage into many shardings like Zilliqa [Tea+17] and Omnidex [KJG+18].
 - **Drawbacks:** Security issues; Incentive mechanism;
- Method 3: **Merging Transactions**
 - Record fewer transactions using payment channels e.g. lightning network [PD16]
 - **Drawbacks:** Only support pairwise payment relationship;

Possible Solutions

- Method 1: **Storage Rent**
 - Miners could know and update ledger state by outsourcing local database to third party or sharing common files with other miners in IPFS system.
 - **Drawbacks:** The operation cost is very high; less competitive; low I/O operations.
- Method 2: **Sharding**
 - Partition the whole blockchain/storage into many shardings like Zilliqa [Tea+17] and Omniledger [KJG+18].
 - **Drawbacks:** Security issues; Incentive mechanism;
- Method 3: **Merging Transactions**
 - Record fewer transactions using payment channels e.g. lightning network [PD16]
 - **Drawbacks:** Only support pairwise payment relationship;
- Method 4: **Compacting Blocks** [Poe16]

Stateless Blockchain

Stateless Blockchain

- This design concept of **Stateless Blockchain** is referred to Peter Todd's blog post [Tod].

Stateless Blockchain

- This design concept of **Stateless Blockchain** is referred to Peter Todd's blog post [Tod].
- Nodes might participate in transaction validation without storing the entire state of the ledger.

Stateless Blockchain

- This design concept of **Stateless Blockchain** is referred to Peter Todd's blog post [Tod].
- Nodes might participate in transaction validation without storing the entire state of the ledger.
- Transaction would include membership proofs for all its input.

Stateless Blockchain

- This design concept of **Stateless Blockchain** is referred to Peter Todd's blog post [Tod].
- Nodes might participate in transaction validation without storing the entire state of the ledger.
- Transaction would include membership proofs for all its input.
- A node would only need to store the current state and verify transactions by checking membership proofs against accumulator state.

Stateless Blockchain

- This design concept of **Stateless Blockchain** is referred to Peter Todd's blog post [Tod].
- Nodes might participate in transaction validation without storing the entire state of the ledger.
- Transaction would include membership proofs for all its input.
- A node would only need to store the current state and verify transactions by checking membership proofs against accumulator state.
- Many optimized schemes are put forward in the forum, such as Merkle Mountain tree [TO] and asynchronous accumulator [RY16] for dual accumulator.

Stateless Blockchain

D. Leung, A. Suhl, Y. Gilad, and N. Zeldovich, “Vault: Fast bootstrapping for the algorand cryptocurrency,” , NDSS, 2019

Vault Techniques


- **Vault:** Reducing the cost of storage and bootstrapping without weakening security guarantees.

Approach	Challenge	Vault's Solution
Reduce state transmitted: Garbage collection	Transaction replay attacks	Force transactions to expire
Reduce state transmitted: Shard state	Small shards lose security	Adaptive Merkle Tree sharding
Reduce size of state proof: Compress history	Attacker tampers with history	Succinct certificates

Vault: Forcing Transactions to Expire


- All transactions contain the fields $t_{issuance}$ and t_{expiry} .
- We define $0 \leq t_{expiry} - t_{issuance} \leq t_{max}$
- The choice of T_{max} affects two considerations.
 - The block number of transactions to detect double spend.
 - Expiration mechanism requires the issuer to reissue expired transactions.
- Support off-chain payment channels e.g. Lightning Network.

Vault: Sharding Balance Storage




- **Shard Witness:** Transactions include Merkle witness for source and destination accounts.

Vault: Sharding Balance Storage



- **Shard Witness:** Transactions include Merkle witness for source and destination accounts.
- **Updating Witness:** The witness could be updated without requiring the issuer to resign the entire transaction.

Vault: Sharding Balance Storage



- **Shard Witness:** Transactions include Merkle witness for source and destination accounts.
- **Updating Witness:** The witness could be updated without requiring the issuer to resign the entire transaction.
- **Adaptive Sharding:** Truncating witness.

Vault: Compressing History

- Skipping Blocks




Fig. 4 : An illustration for skipping blocks

Vault: Compressing History

- Skipping Blocks




Fig. 4 : An illustration for skipping blocks

- Shrinking Certificates

Stateless Blockchain

**A. Chepurnoy, C. Papamanthou, and Y. Zhang, “Edrax: A
cryptocurrency with stateless transaction validation,” Cryptology ePrint
Archive, Report 2018/968, Tech. Rep., 2018**

EDRAX Architecture

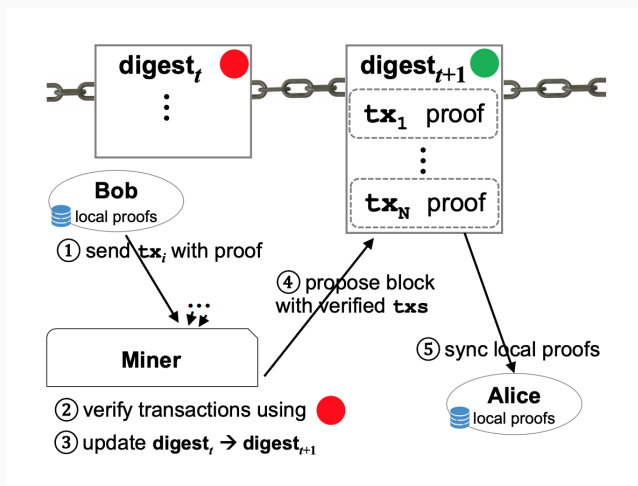


Fig. 5 : Edrax System Architecture

- Represent UTXO as a sparse Merkle tree.

- Represent UTXO as a sparse Merkle tree.
- For account-based model, we can use MHT to represent the mapping of $pk \leftrightarrow \textit{balance}$

- Represent UTXO as a sparse Merkle tree.
- For account-based model, we can use MHT to represent the mapping of $pk \leftrightarrow balance$
- Store the MHT root in the block as digest

- Represent UTXO as a sparse Merkle tree.
- For account-based model, we can use MHT to represent the mapping of $pk \leftrightarrow balance$
- Store the MHT root in the block as digest
- However, during SPEND, we need the Merkle proofs of both sender and receiver

- Represent UTXO as a sparse Merkle tree.
- For account-based model, we can use MHT to represent the mapping of $pk \leftrightarrow balance$
- Store the MHT root in the block as digest
- However, during SPEND, we need the Merkle proofs of both sender and receiver
- Instead, we can use **vector commitment** to represent account-balance mapping.

Edrax in account-based model

- Edrax requires an **one-time setup phase**.
- Each block stores a validation digest
 - A vector commitment digest whose mapping is $i \leftrightarrow h(\text{PK})||\text{balance}$
 - A counter cnt_t which indicates the number of existing accounts
- Each client has a public key in the form $\text{PK} = \text{pk}||i||\text{upk}_i$
- Clients store vector commitment proof π w.r.t. their index i
- The index i is assigned when client submits an **INIT transaction** $[\text{INIT}, \text{pk}]$
 - Miner verifies transaction signature
 - Client is assigned to $i \leftarrow \text{cnt}_t$
 - Miner updates $\text{cnt}_{t+1} \leftarrow \text{cnt}_t + 1$ and a new digest
 $\text{UpdateDigest}(\text{dig}, \text{cnt}, h(\text{PK})||0, \text{upk}_{\text{cnt}})$
 - Clients synchronize proofs accordingly using `UpdateProof`

- To spend δ to client with public key $PK_b = pk_b || j || upk_j$
- Client sends transaction $[PK_a, PK_b, v, \pi_i, v']$ with signature sig
- Miner verifies:
 - sig is valid
 - $v \leq v'$
 - $\text{Verify}(dig_t, i, h(PK_a) || v', \pi_i, vrk)$ passes
- Miner updates digest
 - $dig' \leftarrow \text{UpdateDigest}(dig_t, i, -v, upk_i)$
 - $dig_{t+1} \leftarrow \text{UpdateDigest}(dig', j, v, upk_j)$
- Clients synchronize proofs accordingly using `UpdateProof`

- Either dynamic accumulator or vector commitment provides a **short binding commitment** to a set of elements together with short **membership/non-membership proof**.

- Either dynamic accumulator or vector commitment provides a **short binding commitment** to a set of elements together with short **membership/non-membership proof**.
- In essence, stateless blockchain reduces the storage burden for performing transaction validation by **sacrificing communication overhead**.

- Either dynamic accumulator or vector commitment provides a **short binding commitment** to a set of elements together with short **membership/non-membership proof**.
- In essence, stateless blockchain reduces the storage burden for performing transaction validation by **sacrificing communication overhead**.
- Possible optimizations:

- Either dynamic accumulator or vector commitment provides a **short binding commitment** to a set of elements together with short **membership/non-membership proof**.
- In essence, stateless blockchain reduces the storage burden for performing transaction validation by **sacrificing communication overhead**.
- Possible optimizations:
 - Shorten membership proof size

- Either dynamic accumulator or vector commitment provides a **short binding commitment** to a set of elements together with short **membership/non-membership proof**.
- In essence, stateless blockchain reduces the storage burden for performing transaction validation by **sacrificing communication overhead**.
- Possible optimizations:
 - Shorten membership proof size
 - aggregate membership proofs for a batch of transactions [BBF18]

Possible Works and Directions

- More efficient authenticated data structure instead of sparse merkle tree

Possible Works and Directions

- More efficient authenticated data structure instead of sparse merkle tree
- Local proof synchronization by miner

Possible Works and Directions

- More efficient authenticated data structure instead of sparse merkle tree
- Local proof synchronization by miner
- Introducing proof-serving nodes

Possible Works and Directions

- More efficient authenticated data structure instead of sparse merkle tree
- Local proof synchronization by miner
- Introducing proof-serving nodes
- Supporting smart contracts in the stateless setting

Possible Works and Directions

- More efficient authenticated data structure instead of sparse merkle tree
- Local proof synchronization by miner
- Introducing proof-serving nodes
- Supporting smart contracts in the stateless setting
- Adding privacy to account-based model

References

- [BBF18] D. Boneh, B. Bünz, and B. Fisch, “Batching techniques for accumulators with applications to iops and stateless blockchains,” *Cryptology ePrint Archive*, Report 2018/1188, Tech. Rep, Tech. Rep., 2018.
- [CPZ18] A. Chepurnoy, C. Papamanthou, and Y. Zhang, “Edrax: A cryptocurrency with stateless transaction validation,” *Cryptology ePrint Archive*, Report 2018/968, Tech. Rep., 2018.
- [KJG+18] E. Kokoris-Kogias, P. Jovanovic, L. Gasser, N. Gailly, E. Syta, and B. Ford, “Omniledger: A secure, scale-out, decentralized ledger via sharding,” in *2018 IEEE Symposium on Security and Privacy (SP)*, IEEE, 2018, pp. 583–598.
- [LSGZ19] D. Leung, A. Suhl, Y. Gilad, and N. Zeldovich, “Vault: Fast bootstrapping for the algorand cryptocurrency,” *NDSS*, 2019.
- [PD16] J. Poon and T. Dryja, *The bitcoin lightning network: Scalable off-chain instant payments*, 2016.
- [Poe16] A. Poelstra, *Mimblewimble*, 2016.
- [RY16] L. Reyzin and S. Yakoubov, “Efficient asynchronous accumulators for distributed pki,” in *International Conference on Security and Cryptography for Networks*, Springer, 2016, pp. 292–309.

References

- [Tea+17] Z. Team *et al.*, *The zilliqa technical whitepaper*, 2017.
- [TO] P. Todd and OpenTimestamps, *Merkle mountain tree*, <https://github.com/opentimestamps/opentimestamps-server/blob/master/doc/merkle-mountain-range.md>, Accessed: 2019-04-15.
- [Tod] P. Todd, *Making utxo set growth irrelevant with low- latency delayed txo commitments*, <https://petertodd.org/2016/delayed-txo-commitments>, Accessed: 2019-04-15.
- [Woo+] G. Wood *et al.*, “Ethereum: A secure decentralised generalised transaction ledger,”,

Thanks
Questions?