

# A Preliminary Investigation into Storageless Blockchain

---

Yang Ji

April 15, 2019

1. Introduction
2. Analysis of Existing Works
3. Stateless Blockchain

# Introduction


---

# Background

- In cryptocurrency, peer-to-peer payment transactions are asynchronously broadcasted and recorded in an ordered ledger.
- Consensus protocol requires nodes to validate new transactions.
- e.g. To send 6 ETC from Alice to Bob requires that Alice has at least 6 ETC.
- Simply querying the history on the blockchain is infeasible and insecure due to the large size of blocks.
- For that reason, most cryptocurrency nodes need to locally maintain the **validation state**, which means downloading all past transactions/accounts.

# Background

- In Bitcoin, Zcash and Komodo, **validation state** is a set of immutable coins called **UTXO** (unspent transaction output)



**Fig. 1 : UTXO Model**

# Background

- In Ethereum, NXT and Bitshares organize **validation state** as a set of mutable (and potential long-living) **accounts**.

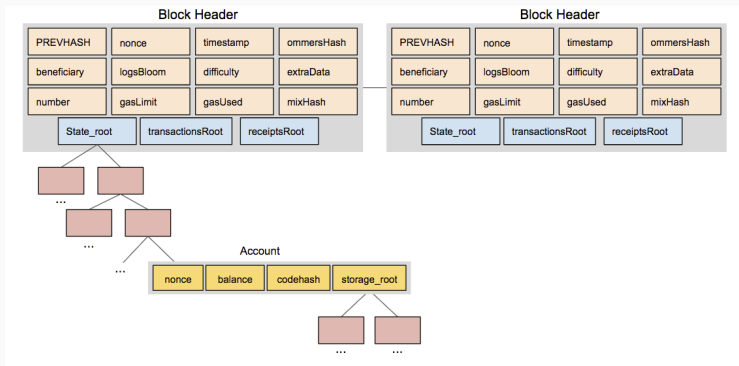


Fig. 2 : Account-based Model

- Other cryptocurrencies

System	Execution State	Proof Size	Bootstrap Security
Bitcoin [21]	UTXOs	Headers + TXs	Probabilistic (heaviest chain wins)
Ethereum [9]	All accounts	Headers + All accounts	Probabilistic (heaviest chain wins)
Permissioned	<u>Live accounts</u> Shards	Majority of trust set's signatures	Cryptographic if majority never compromised; none otherwise
OmniLedger [17] + Chainiac [22]	<u>UTXOs</u> Shards	$\frac{\text{Headers} + \text{Certificates}}{\text{Sparseness}} + \frac{\text{UTXOs}}{\text{Shards}}$	Cryptographic with static attacker; none with adaptive attacker
Algorand [14]	UTXOs	Headers + Certificates + TXs	Cryptographic
Vault	<u>Live accounts</u> Shards	$\frac{\text{Headers} + \text{Certificates}}{\text{Sparseness}} + \frac{\text{Live accounts}}{\text{Shards}}$	Cryptographic

Fig. 3 : Comparison of different cryptocurrency models

- Maintaining ledger state is cumbersome from the perspectives of **storage** and **bootstrapping**
  - Large size (Bitcoin is around 150 GB and Ethereum has exceeded 400 GB)
  - Data storage size is linear with block number and could grow substantially in the coming years
  - Slow disk I/O operations (LevelDB or RocksDB)
  - DoS attack (adversarially-crafted transactions that need massive of disk accesses)
  - Increase the possibility of centralization in blockchains.
- Several developers also talked about **storageless clients** for Bitcoin in 2013.



- State set growth is driven by a number of factors, including the following fact.
  - For **Bitcoin**, there are massive of merge inputs, lost coins and dust outputs.
  - For **Ethereum**, smart contract inadvertently created many **zero-balance accounts** (account for around 38%).
- Based on these, a naïve way is to prune and clean 'useless' dust coins/accounts.
- However, there is little incentive to carry it out for many reasons:
  - Dust coins can't be economically spent and have other use cases.
  - We can't delete zero accounts because nodes need to keep track of the sequence number ("nonce") [Woo+].
- Redesigning index structures and storage mode has limited improvement.

# Analysis of Existing Works

---

# Possible Solutions

- Method 1: Storage Rent

# Possible Solutions

- Method 1: **Storage Rent**
  - Miners could know and update ledger state by outsourcing local database to third party or sharing common files with other miners in IPFS system.

# Possible Solutions

- Method 1: **Storage Rent**
  - Miners could know and update ledger state by outsourcing local database to third party or sharing common files with other miners in IPFS system.
  - **Drawbacks:** The operation cost is very high; less competitive; low I/O operations.

# Possible Solutions

- Method 1: **Storage Rent**
  - Miners could know and update ledger state by outsourcing local database to third party or sharing common files with other miners in IPFS system.
  - **Drawbacks:** The operation cost is very high; less competitive; low I/O operations.
- Method 2: **Sharding**

# Possible Solutions

- Method 1: **Storage Rent**
  - Miners could know and update ledger state by outsourcing local database to third party or sharing common files with other miners in IPFS system.
  - **Drawbacks:** The operation cost is very high; less competitive; low I/O operations.
- Method 2: **Sharding**
  - Partition the whole blockchain/storage into many shardings like Zilliqa [Tea+17] and Omniledger [KJG+18].

# Possible Solutions

- Method 1: **Storage Rent**
  - Miners could know and update ledger state by outsourcing local database to third party or sharing common files with other miners in IPFS system.
  - **Drawbacks:** The operation cost is very high; less competitive; low I/O operations.
- Method 2: **Sharding**
  - Partition the whole blockchain/storage into many shardings like Zilliqa [Tea+17] and Omniledger [KJG+18].
  - **Drawbacks:** Security issues; Incentive mechanism;



# Possible Solutions

- Method 1: **Storage Rent**
  - Miners could know and update ledger state by outsourcing local database to third party or sharing common files with other miners in IPFS system.
  - **Drawbacks:** The operation cost is very high; less competitive; low I/O operations.
- Method 2: **Sharding**
  - Partition the whole blockchain/storage into many shardings like Zilliqa [Tea+17] and Omniledger [KJG+18].
  - **Drawbacks:** Security issues; Incentive mechanism;
- Method 3: **Merging Transactions**

# Possible Solutions

- Method 1: **Storage Rent**
  - Miners could know and update ledger state by outsourcing local database to third party or sharing common files with other miners in IPFS system.
  - **Drawbacks:** The operation cost is very high; less competitive; low I/O operations.
- Method 2: **Sharding**
  - Partition the whole blockchain/storage into many shardings like Zilliqa [Tea+17] and Omniledger [KJG+18].
  - **Drawbacks:** Security issues; Incentive mechanism;
- Method 3: **Merging Transactions**
  - Record fewer transactions using payment channels e.g. lightning network [PD16]

# Possible Solutions

- Method 1: **Storage Rent**
  - Miners could know and update ledger state by outsourcing local database to third party or sharing common files with other miners in IPFS system.
  - **Drawbacks:** The operation cost is very high; less competitive; low I/O operations.
- Method 2: **Sharding**
  - Partition the whole blockchain/storage into many shardings like Zilliqa [Tea+17] and Omniledger [KJG+18].
  - **Drawbacks:** Security issues; Incentive mechanism;
- Method 3: **Merging Transactions**
  - Record fewer transactions using payment channels e.g. lightning network [PD16]
  - **Drawbacks:** Only support pairwise payment relationship;

# Possible Solutions

- Method 1: **Storage Rent**
  - Miners could know and update ledger state by outsourcing local database to third party or sharing common files with other miners in IPFS system.
  - **Drawbacks:** The operation cost is very high; less competitive; low I/O operations.
- Method 2: **Sharding**
  - Partition the whole blockchain/storage into many shardings like Zilliqa [Tea+17] and Omniledger [KJG+18].
  - **Drawbacks:** Security issues; Incentive mechanism;
- Method 3: **Merging Transactions**
  - Record fewer transactions using payment channels e.g. lightning network [PD16]
  - **Drawbacks:** Only support pairwise payment relationship;
- Method 4: **Compacting Blocks** [Poe16]

# Stateless Blockchain

---

# Stateless Blockchain

- This design concept of **Stateless Blockchain** is referred to Peter Todd's blog post [Tod].
- Nodes might participate in transaction validation without storing the entire state of the ledger.
- Transaction would include membership proofs for all its input.
- A node would only need to store the current state and verify transactions by checking membership proofs against accumulator state.
- Many optimized schemes are put forward in the forum, such as TXO commitment [TO] and asynchronous accumulator [RY16] for dual accumulator.

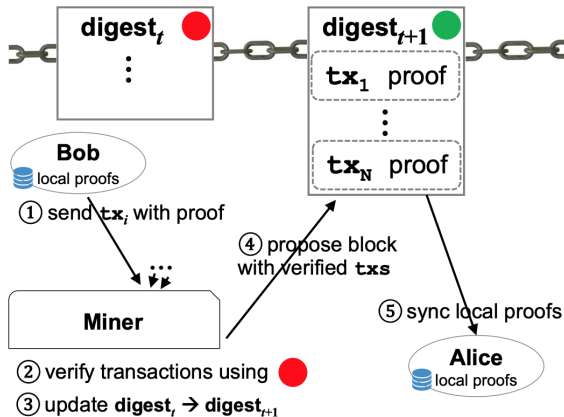


Fig. 4 : System Architecture

## **Stateless Blockchain**

**D. Leung, A. Suhl, Y. Gilad, and N. Zeldovich, “Vault: Fast bootstrapping for the algorand cryptocurrency,” , NDSS, 2019**



# Vault Techniques


- **Vault:** Reducing the cost of storage and bootstrapping without weakening security guarantees.

Approach	Challenge	Vault's Solution
Reduce state transmitted: Garbage collection	Transaction replay attacks	Force transactions to expire
Reduce state transmitted: Shard state	Small shards lose security	Adaptive Merkle Tree sharding
Reduce size of state proof: Compress history	Attacker tampers with history	Succinct certificates

## Vault: Forcing Transactions to Expire

- All transactions contain the fields  $t_{issuance}$  and  $t_{expiry}$ .
- We define  $0 \leq t_{expiry} - t_{issuance} \leq t_{max}$
- The choice of  $T_{max}$  affects two considerations.
  - The block number of transactions to detect double spend.
  - Expiration mechanism requires the issuer to reissue expired transactions.
- Support off-chain payment channels e.g. Lightning Network.


# Vault: Sharding Balance Storage



- **Shard Witness:** Transactions include Merkle witness for source and destination accounts.
- **Updating Witness:** The witness could be updated without requiring the issuer to resign the entire transaction.
- **Adaptive Sharding:** Truncating witness.

# Vault: Compressing History

- Skipping Blocks



**Fig. 5 :** An illustration for skipping blocks

- Shrinking Certificates

## **Stateless Blockchain**

**A. Chepurnoy, C. Papamanthou, and Y. Zhang, “Edrax: A  
cryptocurrency with stateless transaction validation,” Cryptology ePrint  
Archive, Report 2018/968, Tech. Rep., 2018**

- Represent UTXO as a sparse Merkle tree.
- For account-based model, we can use MHT to represent the mapping of  $pk \leftrightarrow balance$
- Store the MHT root in the block as digest
- However, during SPEND, we need the Merkle proofs of both sender and receiver
- Instead, we can use **vector commitment** to represent account-balance mapping.

- To spend  $\delta$  to client with public key  $PK_b = pk_b || j || upk_j$
- Client sends transaction  $[PK_a, PK_b, v, \pi_i, v']$  with signature  $sig$
- Miner verifies:
  - $sig$  is valid
  - $v \leq v'$
  - $\text{Verify}(dig_t, i, h(PK_a) || v', \pi_i, vrk)$  passes
- Miner updates digest
  - $dig' \leftarrow \text{UpdateDigest}(dig_t, i, -v, upk_i)$
  - $dig_{t+1} \leftarrow \text{UpdateDigest}(dig', j, v, upk_j)$
- Clients synchronize proofs accordingly using `UpdateProof`

- Either authenticated data structure or vector commitment provides a short binding commitment to a set of elements together with short membership/non-membership proof.



- Either authenticated data structure or vector commitment provides a short binding commitment to a set of elements together with short membership/non-membership proof.
- In essence, stateless blockchain reduces the storage burden for performing transaction validation by sacrificing communication overhead and assigning miners' tasks to clients.

- Either authenticated data structure or vector commitment provides a **short binding commitment** to a set of elements together with short **membership/non-membership proof**.
- In essence, stateless blockchain reduces the storage burden for performing transaction validation by **sacrificing communication overhead and assigning miners' tasks to clients**.
- Possible optimizations:

- Either authenticated data structure or vector commitment provides a **short binding commitment** to a set of elements together with short **membership/non-membership proof**.
- In essence, stateless blockchain reduces the storage burden for performing transaction validation by **sacrificing communication overhead and assigning miners' tasks to clients**.
- Possible optimizations:
  - Shorten membership proof size

- Either authenticated data structure or vector commitment provides a **short binding commitment** to a set of elements together with short **membership/non-membership proof**.
- In essence, stateless blockchain reduces the storage burden for performing transaction validation by **sacrificing communication overhead and assigning miners' tasks to clients**.
- Possible optimizations:
  - Shorten membership proof size
  - aggregate membership proofs for a batch of transactions [BBF18]

## Possible Works and Directions

- More efficient authenticated data structure instead of sparse merkle tree

## Possible Works and Directions

- More efficient authenticated data structure instead of sparse merkle tree
- Introducing proof-serving nodes

## Possible Works and Directions

- More efficient authenticated data structure instead of sparse merkle tree
- Introducing proof-serving nodes
- Supporting smart contracts in the stateless setting

## Possible Works and Directions

- More efficient authenticated data structure instead of sparse merkle tree
- Introducing proof-serving nodes
- Supporting smart contracts in the stateless setting
- Adding privacy to account-based model



# References

- [BBF18] D. Boneh, B. Bünz, and B. Fisch, “Batching techniques for accumulators with applications to iops and stateless blockchains,” *Cryptology ePrint Archive*, Report 2018/1188, Tech. Rep, Tech. Rep., 2018.
- [CPZ18] A. Chepurnoy, C. Papamanthou, and Y. Zhang, “Edrax: A cryptocurrency with stateless transaction validation,” *Cryptology ePrint Archive*, Report 2018/968, Tech. Rep., 2018.
- [KJG+18] E. Kokoris-Kogias, P. Jovanovic, L. Gasser, N. Gailly, E. Syta, and B. Ford, “Omniledger: A secure, scale-out, decentralized ledger via sharding,” in *2018 IEEE Symposium on Security and Privacy (SP)*, IEEE, 2018, pp. 583–598.
- [LSGZ19] D. Leung, A. Suhl, Y. Gilad, and N. Zeldovich, “Vault: Fast bootstrapping for the algorand cryptocurrency,” *NDSS*, 2019.
- [PD16] J. Poon and T. Dryja, *The bitcoin lightning network: Scalable off-chain instant payments*, 2016.
- [Poe16] A. Poelstra, *Mimblewimble*, 2016.
- [RY16] L. Reyzin and S. Yakoubov, “Efficient asynchronous accumulators for distributed pki,” in *International Conference on Security and Cryptography for Networks*, Springer, 2016, pp. 292–309.

# References

- [Tea+17] Z. Team *et al.*, *The zilliqa technical whitepaper*, 2017.
- [TO] P. Todd and OpenTimestamps, *Merkle mountain tree*, <https://github.com/opentimestamps/opentimestamps-server/blob/master/doc/merkle-mountain-range.md>, Accessed: 2019-04-15.
- [Tod] P. Todd, *Making utxo set growth irrelevant with low- latency delayed txo commitments*, <https://petertodd.org/2016/delayed-txo-commitments>, Accessed: 2019-04-15.
- [Woo+] G. Wood *et al.*, “Ethereum: A secure decentralised generalised transaction ledger,”,

**Thanks**  
**Questions?**